

Instrumenting Runtime Enforcement

No Institute Given

Abstract. In runtime enforcement, a monitor observes a system’s actions and modifies them to ensure compliance with a specification. In practice, the specification enforced by the monitor is often defined in terms of high-level, abstract events, while the system’s behavior is comprised of low-level, concrete actions. The connection between these two levels is established through an *instrumentation* process where developers must ensure (i) that system actions report the right events to the monitor and (ii) that the modifications of the system’s behavior required by the monitor are correctly enforced. However, the abstraction gap between high-level policies and low-level actions makes this process error-prone. In this paper, we refine the standard formal model of runtime enforcement, in which instrumentation is left implicit, into a more precise model that explicitly accounts for instrumentation. We investigate the correctness of instrumentation in this model and present a novel library, called INSTRLIB, that allows developers to easily instrument Python applications for runtime enforcement.

Keywords: Runtime Enforcement, Instrumentation, Temporal Logic

1 Introduction

In 2022, personal data of roughly a third of Australia’s population was stolen from the telecommunication provider Optus [23]. The attack was unsophisticated, involving the attacker exploiting a coding error in the instrumentation of an internet-facing, legacy API. Due to this error, the access control (AC) mechanism, while in place, was not invoked to protect the legacy API, allowing for the easy retrieval of millions of user records.

This data breach highlights a recurring theme: even when appropriate security mechanisms are in place, their incorrect instrumentation allows attackers to bypass them entirely. A rigorous approach to instrumentation is particularly crucial in applications where the specification is a property of sequences of abstract events, with each event reflecting many possible concrete system actions. This issue materializes particularly clearly when enforcing requirements derived from privacy laws [15]: if a regulation requires that “no user data is used without prior consent,” then being able to ensure that “data usage” is blocked whenever “consent” has not been previously registered is not enough to certify that an application complies with the law. The developers also need to ensure that “data usage” is correctly intercepted by existing system instrumentation whenever low-level actions such as database reads and writes occur; furthermore, they must check that “consent” is only registered when users actually give consent in the UI.




Fig. 1: The “classical” (top) and “extended” (bottom) enforcement models

Runtime enforcement generalizes AC by using execution monitors that observe the actions of a running system and modify these actions to ensure that only compliant behaviors are allowed. While, in recent years, increasingly powerful enforcement approaches and tools have been developed, much less attention has been devoted to how to properly instrument systems or audit existing instrumented systems to ensure correct enforcement.

Figure 1 (top) shows the idealized system model used in most previous work, which we call the “classical” model of runtime enforcement. In this model, the System under Enforcement (SuE) is a labeled transition system (LTS) with transitions of the form $s \xrightarrow{e} s'$ labeled by some *event* e . When attempting to take a transition labeled by e (Step 1), the SuE sends the event e to a policy decision point (PDP) in charge of ensuring the SuE’s compliance with a property P (Step 2). The PDP may edit [18] the event e to a sequence of events e' compliant with P (Step 3). The modification can involve removing, replacing, or inserting events. The events e' are then returned as a command to the SuE (Step 4), which takes the appropriate transitions (Step 5). As it assumes that all events are correctly sent and that the SuE always follows the PDP’s commands, the classical model cannot capture non-compliance due to incorrect instrumentation.

We propose an “extended” model, shown in Figure 1 (bottom) where SuE is modeled as an LTS with transitions $s \xrightarrow{a} s'$ labeled by *actions* that are distinct from the events sent to the PDP. In addition to SuE and PDP, our model introduces an explicit policy enforcement point (PEP) [13] that instruments the actions performed by the SuE to events processed by the PDP. Concretely, the PEP intercepts every action a attempted by the system (Step 2) and maps it to a sequence of events e (Step 3). The sequence e is then sent to the PDP (Step 4), which edits e to some e' (Step 5) and returns e' (Step 6). The PEP maps e' back to a sequence of actions a' (Step 7) that PEP enforces in the SuE (Steps 8 and 9).

In our extended model, the SuE’s specification consists of two elements: the property P and a *mediator* I providing the desired interpretation of system actions in terms of PDP events. In the Optus data breach, the property P may have been “whenever a user u is performs an API access (`APIAccess( $u$ )` event), then u is authenticated”. The mediator I should have mapped both system actions `legacyAPIAccess( $u$ )` and `modernAPIAccess( $u$ )` to the event `APIAccess( $u$ )`. However, the instrumentation step (2) failed to map `LegacyAPIAccess( $u$ )` to

$\text{APIAccess}(u)$. In this case, despite the PDP correctly enforcing P , the composition of the SuE, PEP, and PDP did not provide the desired security guarantees.

Contributions. After reviewing the classical enforcement model found in previous work (Section 2), we make the following contributions:

- We formally introduce our extended enforcement model and define appropriate notions of PEP correctness which provide necessary and sufficient conditions for the correctness of the instrumentation (Section 3).
- We show how these conditions can be weakened for properties for which the occurrence of certain events can be soundly overapproximated (Section 4).
- We specialize our theory to a framework for instrumenting applications using a family of state-of-the-art PDPs which process events in finite sets (‘batches’). We provide a correct PEP algorithm for this setup and give auditing check-lists that can be used to guarantee compliance (Section 5).
- We implement our algorithm in INSTRLIB, an open-source library for enforcing properties of Python applications using the state-of-the-art ENFGUARD [17] tool. We illustrate our framework in a case study, implementing privacy requirements in a micro-blogging application using INSTRLIB and then auditing our instrumentation’s correctness (Section 6).

Related Work. Models of runtime enforcement mechanisms include security automata [11,24], edit automata [18], mandatory results automata (MRA) [19], and timed automata [3,21]. More recently, several frameworks for enforcing expressive first-order policies at runtime have also been developed [14,16,17]. Runtime enforcement can be performed both on high-level events and low-level system actions, e.g., through inlining [10]. Most models only consider the PDP’s logic, without any guarantees of correct instrumentation. MRAs distinguish between system *actions* and enforced *results* and may fulfill the role of both a PDP and PEP, but do not provide for a clear separation between instrumentation and property enforcement. In contrast, several works from the security community discuss the composition of a PDP and PEP in the context of runtime enforcement while lacking a rigorous theory to describe this composition [4,22,15]. Our account of the ‘classical model’ builds on work by Aceto et al. [1] and Falcone et al. [8,12], where an SuE modeled as an LTS and the PDP run in lockstep.

2 Preliminaries

After introducing notation, we review labeled transition systems and edit automata (Section 2.1). We then introduce the ‘classical’ enforcement model and its associated notions of PDP soundness and transparency (Section 2.2).

Notation. Given a set A , the set of all finite sequences of elements of A is denoted by A^* . We use Greek letters (e.g., $\alpha, \sigma, \rho, \dots$) or bold Latin letters (e.g., $\mathbf{a}, \mathbf{e}, \mathbf{\ell}, \dots$) to denote sequences; bold Greek letters denote sequences of sequences. We denote the empty sequence as ε and finite sequences as $\langle a_1, a_2, \dots, a_n \rangle$. The sequence $\sigma \setminus a$ stands for σ with all occurrences of a removed, $\sigma_{..i}$ for the prefix of σ of length i , and $\text{pre}(\sigma)$ for the set of all prefixes of σ . Given $\sigma, \sigma' \in A^*$, the

sequence $\sigma \cdot \sigma'$ is the concatenation of σ and σ' . Given a sequence of sequences $\boldsymbol{\sigma} = \langle \sigma_1, \dots, \sigma_n \rangle \in (A^*)^*$, we denote by $\circ \boldsymbol{\sigma} \triangleq \sigma_1 \cdot \sigma_2 \cdot \dots \cdot \sigma_n$ the concatenation of the σ_i .

For any set of labels $\mathcal{L}$, the set of *traces* over $\mathcal{L}$ is $\mathbb{T}_{\mathcal{L}} \triangleq \mathcal{L}^*$. A *property* $P_{\mathcal{L}}$ over $\mathcal{L}$ is a set of traces, i.e., a subset $P_{\mathcal{L}} \subseteq \mathbb{T}_{\mathcal{L}}$. To support *internal* (or ‘silent’) system actions, we use a distinguished label $\tau \notin \mathcal{L}$ and denote $\mathcal{L}_{\tau} \triangleq \mathcal{L} \cup \{\tau\}$.

2.1 Labeled Transition Systems and Edit Automata

As in previous work [1,7,2], we model systems as labeled transition systems:

Definition 1 (LTS). A labeled transition system (LTS) over $\mathcal{L}$ is a quadruple $\mathcal{S} = (\mathbb{S}, \mathcal{L}, s^0, \dot{\rightarrow})$ such that $\mathbb{S}$ is a set of states, $\mathcal{L}$ is a set of labels, $s^0 \in \mathbb{S}$ is an initial state, and $\dot{\rightarrow} \subseteq \mathbb{S} \times \mathcal{L}_{\tau} \times \mathbb{S}$ is a transition relation labeled by $\mathcal{L}_{\tau}$.

We write $s \xrightarrow{\ell} s'$ for $(s, \ell, s') \in \dot{\rightarrow}$. An LTS execution is of the form $s^0 \xrightarrow{\ell_1} s^1 \xrightarrow{\ell_2} \dots \xrightarrow{\ell_n} s^n$ and the trace of such an execution is $\sigma = \langle \ell_1, \ell_2, \dots \rangle \setminus \tau$, removing all internal τ actions. In this case, we also write $s^0 \xrightarrow{\sigma} s^n$.

Example 1. Figure 2 represents a social network application as an LTS over two different sets of labels, $\mathcal{E}$ (‘events’, left) and $\mathcal{A}$ (‘actions’, right) capturing the system’s behavior at two levels of abstraction. For simplicity, we model the system for a single user. At a high level (events), the system can be described in terms of user interaction (give **Consent** for data usage, **Revoke** consent, **Request** data deletion), backend operations (**Use**, **Delete** user data), and clock ticks (**Tick**). At a lower level (actions), one can observe UI interactions (**click_yes**, **click_no** in a consent banner, clicks on a **request** button), database (**read**, **write**, **delete**) or authentication (**login**, **logout**) operations, and clock ticks (**tick**).

Edit automata (EA) [18] are a general model for PDPs, providing an abstract model for a large class of practical enforcement mechanisms. Edit automata are a special kind of LTS which, in each step, reads a label ℓ^1 from some set of labels $\mathcal{L}^1$ and edits it deterministically into a possibly empty sequence of labels ℓ_2 from another set $\mathcal{L}^2$. If no ℓ_2 exists for a given ℓ_1 , the processing of the trace is blocked, similar to execution cutting in security automata [24].

Definition 2 (Edit Automaton). An edit automaton (EA) over $(\mathcal{L}^1, \mathcal{L}^2)$ is an LTS $\delta = (\mathbb{S}_{\delta}, \mathcal{L}^1 \times \mathcal{L}_{\tau}^{2*}, s_{\delta}^0, \xrightarrow{\triangleright}_{\delta})$ with a transition relation labeled with pairs of labels $(\ell_1, \ell_2) \in \mathcal{L}^1 \times \mathcal{L}_{\tau}^{2*}$ also denoted as $\ell_1 \triangleright \ell_2$, such that, for any $s_{\delta} \in \mathbb{S}_{\delta}$ and $\ell \in \mathcal{L}_1$, there exists at most one pair $(s'_{\delta}, \ell_2) \in \mathbb{S}_{\delta} \times \mathcal{L}_{\tau}^{2*}$ such that $s_{\delta} \xrightarrow{\ell \triangleright \ell_2} s'_{\delta}$.

An edit automaton is input-enabled [20] iff for any ℓ^1 , the automaton can always edit ℓ^1 to some ℓ_2 .

Definition 3. An edit automaton δ over $(\mathcal{L}^1, \mathcal{L}^2)$ is input-enabled iff for any $s_{\delta} \in \mathbb{S}_{\delta}$ and $\ell_1 \in \mathcal{L}^1$, there exists $s'_{\delta} \in \mathbb{S}_{\delta}$ and $\ell_2 \in \mathcal{L}_{\tau}^{2*}$ such that $s_{\delta} \xrightarrow{\ell_1 \triangleright \ell_2} s'_{\delta}$.




Fig. 2: Systems under Enforcement (SuE) in the Classical and Extended Model




Fig. 3: Policy Decision Points (PDPs) and Policy Enforcement Point (PEP)

Note that there are two interpretations of an EA: the EA accepts a language of traces over pairs containing a label (from $\mathcal{L}^1$) and a sequence of labels (from $\mathcal{L}_\tau^{2*}$). Alternatively, it is a transducer, translating a sequence of labels from $\mathcal{L}^1$ into a sequence of labels from $\mathcal{L}^2$ by concatenating the ℓ_2 . For $n \in \mathbb{N}$, $\xi = (x_i)_{1 \leq i < n}$, and $v = (y_i)_{1 \leq i < n}$, we denote by $\xi \triangleright v$ the zipped sequence $(x_i \triangleright y_i)_{1 \leq i < n}$. For any EA δ , we abuse notation to use δ as a partial function such that $\delta(\sigma) = \circ \sigma' \iff \exists s'_\delta. s_\delta^0 \xrightarrow{\sigma \triangleright \sigma'} s'_\delta$, reflecting the view of δ as a trace transducer.

2.2 The Classical Model of Runtime Enforcement

In the classical enforcement model [18,1,7], an SuE (LTS) and a PDP (edit automaton) over the same set of labels are composed, progressing in lock-step. This can be formalized as follows in terms of LTS composition:

Definition 4. An LTS $\mathcal{S}$ over $\mathcal{L}$ and an edit automaton δ over $(\mathcal{L}, \mathcal{L})$ serving as a PDP can be composed into an LTS $\langle \mathcal{S} | \delta \rangle = (\mathbb{S}_{\langle \mathcal{S} | \delta \rangle}, \mathcal{L}, s_{\langle \mathcal{S} | \delta \rangle}^0, \rightarrow_{\langle \mathcal{S} | \delta \rangle})$ by

$$\begin{aligned} \mathbb{S}_{\langle \mathcal{S} | \delta \rangle} &\triangleq \mathbb{S}_{\mathcal{S}} \times \mathbb{S}_{\delta} \times \mathcal{L}^* & s_{\langle \mathcal{S} | \delta \rangle}^0 &\triangleq (s_{\mathcal{S}}^0, s_{\delta}^0, \varepsilon) \\ \frac{s_{\mathcal{S}} \xrightarrow{\tau} s'_{\mathcal{S}}}{(s_{\mathcal{S}}, s_{\delta}, b) \xrightarrow{\tau}_{\langle \mathcal{S} | \delta \rangle} (s'_{\mathcal{S}}, s_{\delta}, b)} \text{step}_{\tau} & \quad \frac{s_{\mathcal{S}} \xrightarrow{\ell' \neq \tau} s''_{\mathcal{S}} \quad s_{\delta} \xrightarrow{\ell' \triangleright \varepsilon}_{\delta} s'_{\delta}}{(s_{\mathcal{S}}, s_{\delta}, \varepsilon) \xrightarrow{\tau}_{\langle \mathcal{S} | \delta \rangle} (s_{\mathcal{S}}, s'_{\delta}, \varepsilon)} \text{step}_0 \\ \frac{s_{\mathcal{S}} \xrightarrow{\ell' \neq \tau} s''_{\mathcal{S}} \quad s_{\delta} \xrightarrow{\ell' \triangleright (\langle \ell \rangle \cdot \ell)}_{\delta} s'_{\delta} \quad s_{\mathcal{S}} \xrightarrow{\ell} s'_{\mathcal{S}}}{(s_{\mathcal{S}}, s_{\delta}, \varepsilon) \xrightarrow{\ell}_{\langle \mathcal{S} | \delta \rangle} (s'_{\mathcal{S}}, s'_{\delta}, \ell)} \text{step}_{+} & \quad \frac{s_{\mathcal{S}} \xrightarrow{\ell} s'_{\mathcal{S}}}{(s_{\mathcal{S}}, s_{\delta}, \langle \ell \rangle \cdot \ell) \xrightarrow{\ell}_{\langle \mathcal{S} | \delta \rangle} (s'_{\mathcal{S}}, s_{\delta}, \ell)} \text{buf} \end{aligned}$$

The states of the *enforced LTS* $\langle \mathcal{S} | \delta \rangle$ have three components: the state of $\mathcal{S}$, the state of δ , and a buffer in $\mathcal{L}^*$ used to temporarily store the sequences of labels returned by the PDP δ until corresponding transitions are taken by the system. The enforced LTS has four kinds of transitions:

- step_{τ} transitions mirror internal transitions in $\langle \mathcal{S} | \delta \rangle$ without δ 's intervention;
- step_0 transitions are triggered by the EA erasing the original label ℓ' of a transition of $\mathcal{S}$, i.e., editing it to ε . In this case, the substate of $\langle \mathcal{S} | \delta \rangle$ corresponding to the state of $\mathcal{S}$ does not change;
- step_{+} transitions are triggered by the EA editing the original label ℓ' of a transition of $\mathcal{S}$ to a non-empty sequence of labels $\langle \ell \rangle \cdot \ell$. In this case, a first transition labeled by ℓ is performed immediately in $\mathcal{S}$, while ℓ is buffered;
- buf transitions are performed while the buffer of the enforced LTS is not empty (step_0 and step_{+} are disallowed in this case), performing one transition labeled with the first label in the buffer and removing it from the buffer.

This model supports the interruption of system execution by the PDP: whenever the system can take a step $s_{\mathcal{S}} \xrightarrow{\ell'} s'_{\mathcal{S}}$ but there is no ℓ such that $s_{\delta} \xrightarrow{\ell} s'_{\delta}$, the execution is blocked.

Example 2. We compose the SuE over $\mathcal{E}$ in Figure 2 with the PDPs P_{Consent} and P_{Deletion} in Figure 3. Consider our LTS on $\mathcal{E}$ in state *Timeline* (the user is logged in) running together the PDP P_{Consent} in state *RVK* (i.e., user consent was revoked or never given). If the SuE attempts *Use*, the PDP receives *Use* in *RVK* and suppresses it by taking the step $RVK \xrightarrow{\text{Use} \triangleright \varepsilon} RVK$; by rule step_0 , the state of the composed LTS is left unchanged. Now, compose our LTS on $\mathcal{E}$ in the state *Timeline* with the PDP P_{Delete} . Assume that the LTS just executed *Request*, which put the PDP in state *Wait*. If the SuE attempts *Tick*, the PDP takes a step $Wait \xrightarrow{\text{Tick} \triangleright (\text{Delete}, \text{Tick})}_{\delta} Base$, inserting *Delete*; the SuE executes *Timeline* $\xrightarrow{\text{Delete}} Timeline$ and buffers $\langle \text{Tick} \rangle$ by rule step_{+} ; finally, since the buffer is not empty, the SuE takes a transition *Timeline* $\xrightarrow{\text{Tick}} Timeline$ by rule buf .

A PDP's correctness is generally expressed in terms of two standard notions [18]: soundness and transparency. Soundness with respect to some property P states that any prefix of a trace edited by the PDP is in P ; transparency states that traces that already adhere to P are not modified by the PDP.

Definition 5 (Soundness). An edit automaton δ over $(\mathcal{L}, \mathcal{L})$ is sound with respect to a property $P \subseteq \mathbb{T}_{\mathcal{L}}$ iff for any $\sigma \in \mathbb{T}_{\mathcal{L}}$, $\text{pre}(\delta(\sigma)) \subseteq P$.

Definition 6 (Transparency). An edit automaton δ over $(\mathcal{L}, \mathcal{L})$ is transparent with respect to a property $P \subseteq \mathbb{T}_{\mathcal{L}}$ iff for any $\sigma \in P$, $\delta(\sigma) = \{\sigma\}$.

Runtime enforcement aims to ensure the compliance of the enforced LTS with some P . Another common requirement is that, if an execution of the original LTS already fulfills P , then this execution is not altered by the enforcement mechanism. The corresponding properties of the *enforced system* are as follows:

Definition 7 (Compliance). An LTS $\mathcal{S} = (\mathbb{S}, \mathcal{L}, s^0, \dot{\rightarrow})$ complies with the property P iff for any execution $s^0 \xrightarrow{\sigma} s$, we have $\sigma \in P_{\mathcal{L}}$.

Definition 8 (Preservation). An LTS $\mathcal{S} = (\mathbb{S}, \mathcal{L}, s^0, \dot{\rightarrow})$ preserves the property $P_{\mathcal{L}}$ with respect to an LTS $\mathcal{S}_{\star} = (\mathbb{S}_{\star}, \mathcal{L}, s_{\star}^0, \dot{\rightarrow}_{\star})$ iff for any execution $s_{\star}^0 \xrightarrow{\sigma}_{\star} s_{\star}$ of $\mathcal{S}_{\star}$ such that $\sigma \in P_{\mathcal{L}}$, we have an execution $s^0 \xrightarrow{\sigma} s$ of $\mathcal{S}$.

The following theorems state the standard equivalence result for PDPs [18]: a PDP δ is sound iff any enforced system $\langle \mathcal{S} | \delta \rangle$ complies with P ; it is transparent iff any enforced system $\langle \mathcal{S} | \delta \rangle$ preserves P with respect to $\mathcal{S}$.

Theorem 1. The edit automaton δ is sound with respect to property P iff for any LTS $\mathcal{S}$, the LTS $\langle \mathcal{S} | \delta \rangle$ complies with P .

Theorem 2. The edit automaton δ is transparent with respect to property P iff for any LTS $\mathcal{S}$, the LTS $\langle \mathcal{S} | \delta \rangle$ preserves P with respect to $\mathcal{S}$.

3 The Extended Enforcement Model

We now extend the classical model just presented to a model that explicitly distinguishes between (low-level) system actions and high-level events. To this end, we first define a formal notion of a PEP that mediates between these two levels and introduce new notions of PEP soundness and transparency (Section 3.1); then, we describe the three-way composition of an SuE, PEP, and PDP and provide analogues of Theorems 1 and 2 in this extended setup (Section 3.2).

3.1 Formal Model of the Policy Enforcement Point (PEP)

Consider a set of system actions $\mathcal{A}$ and set of events $\mathcal{E}$. A PEP η is a pair of edit automata, namely the *instrumentor* η_i and the *enforcer* η_e , which perform the Steps 2 and 6 in Figure 1 (bottom). The *instrumentor* η_i maps system actions to sequences of PDP events, while the *enforcer* η_e maps PDP events back to sequences of system actions. For the PEP to serve as a transducer between traces of actions and traces of events as they occur in the enforced system, (i) the PEP's state after Step 6 may depend on the edited events it received in Step 5, but not on the attempted actions it had mapped in Step 1. Moreover, (ii) if the sequence of edited actions at the end of Step 6 is empty, then the PEP's state must remain the same as before Step 1. Formalizing this results in the following definition:

Definition 9. A PEP over $(\mathcal{A}, \mathcal{E})$ is a pair $\eta = (\eta_i, \eta_e)$ where $\eta_i = (\mathbb{S}_\eta, \mathcal{A} \times \mathcal{E}_\tau^*, s_\eta^0, \xrightarrow{\triangleright}_{\eta,i})$ is an edit automaton over $(\mathcal{A}, \mathcal{E})$ and $\eta_e = (\mathbb{S}_\eta, \mathcal{E} \times \mathcal{A}_\tau^*, s_\eta^0, \xrightarrow{\triangleright}_{\eta,e})$ is an edit automaton over $(\mathcal{E}, \mathcal{A})$, sharing the same state space, such that for any

$$s_\eta^1 \xrightarrow{a'_1 \triangleright e'_1}_{\eta,i} s_\eta^{2,1} \xrightarrow{e_1 \triangleright \alpha_1}_{\eta,e} s_\eta^{3,1} \quad s_\eta^1 \xrightarrow{a'_2 \triangleright e'_2}_{\eta,i} s_\eta^{2,2} \xrightarrow{e_2 \triangleright \alpha_2}_{\eta,e} s_\eta^{3,2},$$

with $\circ \alpha_1 = \circ \alpha_2$, then (i) $s_\eta^{3,1} = s_\eta^{3,2}$, and (ii) if $\circ \alpha_1 = \varepsilon$, then $s_\eta^{3,1} = s_\eta^1$.

Straightforwardly, we define:

Definition 10. A PEP η is input-enabled iff both η_i and η_e are input-enabled.

Whenever we have a sequence of alternating η_i and η_e transitions $s_\eta^0 \xrightarrow{a'_1 \triangleright e'_1}_{\eta,i} s_\eta^1 \xrightarrow{\sigma_1 \triangleright \rho_1}_{\eta,e} s_\eta^{1,1} \xrightarrow{a'_2 \triangleright e'_2}_{\eta,i} s_\eta^2 \xrightarrow{\sigma_2 \triangleright \rho_2}_{\eta,e} \dots \rightarrow s_\eta^n$, we denote by $s_\eta^0 \xrightarrow{\sigma \triangleright \rho_1 \dots \rho_n}_{\eta} s_\eta^n$ the sequence of transitions generating the action trace $\circ(\rho_1 \dots \rho_n)$ and event trace $\circ \sigma$.

Example 3. The PEP in Figure 3 mediates between $\mathcal{A}$ and $\mathcal{E}$, mapping the intercepted SuE actions into sequences of PDP events and events edited by the PDP back into SuE actions. We want to emphasize two aspects of the PEP mediation, (1) the PEP can not sent events that are irrelevant for the PDP such as **login** and **logout**, (2) the PEP can encode non-injective mappings from actions to events such as **Use** to **Read/Write** through different states.

3.2 A More Realistic Enforcement Model

Having formalized the PEP model, we now compose it with an LTS and PDP:

Definition 11. An LTS $\mathcal{S}$ over $\mathcal{A}$, a PEP η over $(\mathcal{A}, \mathcal{E})$, and a PDP δ over $\mathcal{E}$ can be composed into an LTS $\langle \mathcal{S} | \eta | \delta \rangle = (\mathbb{S}_{\langle \mathcal{S} | \eta | \delta \rangle}, \mathcal{L}, s_{\langle \mathcal{S} | \eta | \delta \rangle}^0, \xrightarrow{\cdot}_{\langle \mathcal{S} | \eta | \delta \rangle})$ by

$$\begin{aligned} \mathbb{S}_{\langle \mathcal{S} | \eta | \delta \rangle} &\triangleq \mathbb{S}_\mathcal{S} \times \mathbb{S}_\eta \times \mathbb{S}_\delta \times \mathcal{L}^* & s_{\langle \mathcal{S} | \eta | \delta \rangle}^0 &\triangleq (s_\mathcal{S}^0, s_\eta^0, s_\delta^0, \varepsilon) \\ & & s_\mathcal{S} &\xrightarrow{a' \neq \tau} s_\mathcal{S}'' \quad s_\eta \xrightarrow{a' \triangleright e'}_{\eta,i} s_\eta'' \\ & & s_\delta &\xrightarrow{e' \setminus \tau \triangleright e}_{\delta} s_\delta' \\ & & s_\eta'' &\xrightarrow{e' \setminus \tau \triangleright e}_{\eta,e} s_\eta' \quad e' \setminus \tau \neq \varepsilon \\ \hline (s_\mathcal{S}, s_\eta, s_\delta, b) &\xrightarrow{\tau}_{\langle \mathcal{S} | \eta | \delta \rangle} (s_\mathcal{S}', s_\eta, s_\delta, b) & \text{step}_\tau & \quad (s_\mathcal{S}, s_\eta, s_\delta, \varepsilon) \xrightarrow{\tau}_{\langle \mathcal{S} | \eta | \delta \rangle} (s_\mathcal{S}, s_\eta', s_\delta', \varepsilon) \quad \text{step}_0 \\ & & & \\ & & s_\mathcal{S} &\xrightarrow{a' \neq \tau} s_\mathcal{S}'' \quad s_\eta \xrightarrow{a' \triangleright e'}_{\eta,i} s_\eta'' \\ & & s_\delta &\xrightarrow{e' \setminus \tau \triangleright e}_{\delta} s_\delta' \quad s_\eta'' \xrightarrow{e' \setminus \tau \triangleright \langle \langle a \rangle \cdot \mathbf{a} \rangle}_{\eta,e} s_\eta' \\ & & s_\mathcal{S} &\xrightarrow{\mathbf{a}} s_\mathcal{S}' \quad e' \setminus \tau \neq \varepsilon \\ \hline (s_\mathcal{S}, s_\eta, s_\delta, \varepsilon) &\xrightarrow{\mathbf{a}}_{\langle \mathcal{S} | \eta | \delta \rangle} (s_\mathcal{S}', s_\eta', s_\delta', \mathbf{a}) & \text{step}_+ & \quad (s_\mathcal{S}, s_\eta, s_\delta, \varepsilon) \xrightarrow{\tau}_{\langle \mathcal{S} | \eta | \delta \rangle} (s_\mathcal{S}', s_\eta, s_\delta, \varepsilon) \quad \text{step}_\varepsilon \\ & & & \\ & & s_\mathcal{S} &\xrightarrow{\mathbf{a}} s_\mathcal{S}' \\ \hline (s_\mathcal{S}, s_\eta, s_\delta, \langle a \rangle \cdot \mathbf{a}) &\xrightarrow{\mathbf{a}}_{\langle \mathcal{S} | \eta | \delta \rangle} (s_\mathcal{S}', s_\eta, s_\delta, \mathbf{a}) & \text{buf} & \end{aligned}$$

The states of the enforced LTS $\langle \mathcal{S} | \eta | \delta \rangle$ have four components, now also including the state of the PEP η . Rules $step_\tau$ and buf are similar to before, while $step_0$ and $step_+$ incorporate the mediation through η_i and η_e . We add a fifth rule $step_\varepsilon$ which covers the case in which the instrumentor generates an empty sequence of events in response to an action, leading the enforced system to take the attempted transition without any intervention from δ .

Example 4. Consider composing our LTS on $\mathcal{A}$ with the PDP P_{Delete} and the PEP in Figure 3. Assume that the SuE is in state *Timeline* and just executed a transition labeled by **request**, putting the PDP in state *Wait*. The PEP is in state *Base*. If the SuE attempts **tick**, the PEP takes a step $Base \xrightarrow{\text{tick} \triangleright \langle \text{Tick} \rangle} \eta_i$; *Base*; the PDP takes a step $Wait \xrightarrow{\text{Tick} \triangleright \langle \text{Delete}, \text{Tick} \rangle} \delta$ *Base*, inserting **Delete**; the PEP takes steps $Base \xrightarrow{\text{Delete} \triangleright \text{delete}} \eta_e$ *Base* $\xrightarrow{\text{Tick} \triangleright \text{tick}} \eta_e$ *Base*; the SuE executes $Timeline \xrightarrow{\text{delete}} Timeline$ and buffers $\langle \text{tick} \rangle$ by rule $step_+$; since the buffer is non-empty, the SuE takes a transition $Timeline \xrightarrow{\text{tick}} Timeline$ by rule buf .

Now, compose our LTS with the PDP P_{Consent} and the PEP as above, with the PDP in state *RVK*. If, in state *Timeline*, the SuE attempts **write**, then the PEP maps $Base \xrightarrow{\text{write} \triangleright \text{Use}} Write$; the PDP receives **Use** without prior consent and suppresses it, taking a step $RVK \xrightarrow{\text{Use} \triangleright \varepsilon} RVK$; consequently, according to rule $step_0$, the SuE does not take any transition.

In our extended model, the PEP is in charge of mediating between traces of actions generated by the system and traces of events observed by the PDP. Its correctness can therefore only be assessed with regard to a mapping between these two sets of traces, which is part of the system's specification and which we call a *trace mediator*. In the following, we adopt the following definition of soundness, which we will show next to provide a necessary and sufficient condition for the compliance of the three-way composition:

Definition 12. A PEP η over $(\mathcal{A}, \mathcal{E})$ is sound with respect to a trace mediator $I : \mathcal{A}^* \rightarrow \mathcal{E}^*$ iff whenever $s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s'_\eta \xrightarrow{a' \triangleright e'} \eta_i$ $s''_\eta \xrightarrow{e \triangleright \alpha} \eta_e$ s'''_η and $\circ \rho \cdot \circ \alpha \in \text{dom}(I)$, then $I(\circ \rho \cdot \circ \alpha)$ is a prefix of $\circ \sigma \cdot e$.

Note that the above definition only requires the action trace $I(\circ \rho \cdot \circ \alpha)$ to be a prefix of the event trace $\circ \sigma \cdot e$, rather than equal to it. Intuitively, we observe that, if the mapped action trace $I(\rho)$ is always a prefix of the event trace ρ , then the fact that $\text{pre}(\sigma) \subseteq P$ guarantees $I(\rho) \in P$.

Similarly, for transparency, we require that, whenever the edit automaton serving as a PDP does not modify the sequence of events in Step 4, the enforcer returns the same action that was originally passed to the instrumentor.

Definition 13. A PEP η is transparent iff it is input-enabled and, whenever $s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s'_\eta \xrightarrow{a' \triangleright e} \eta_i$ $s''_\eta \xrightarrow{e \triangleright \alpha} \eta_e$ s'''_η , then $\circ \alpha = \langle a' \rangle$.

The following theorems provide analogues for the standard correctness results in Theorems 1–2 for our three-way composition. They show that an edit automaton δ is sound (resp. transparent) with respect to a property P if and only if when composed with a sound (resp. transparent) PEP η with respect to a trace mediator I and an arbitrary LTS, the resulting enforced LTS complies with $I^{-1}(P)$ (resp. preserves $I^{-1}(P)$).

Theorem 3. *An edit automaton δ over $\mathcal{E}$ is sound with respect to P iff for any PEP η over $(\mathcal{A}, \mathcal{E})$ that is sound with respect to I and LTS $\mathcal{S}$ over $\mathcal{A}$, the LTS $\langle \mathcal{S} | \eta | \delta \rangle$ complies with $I^{-1}(P) \triangleq \{\rho \in \text{dom}(I) \mid I(\rho) \in P\}$.*

Theorem 4. *An edit automaton δ over $\mathcal{E}$ is transparent with respect to P iff for any transparent PEP η over $(\mathcal{A}, \mathcal{E})$ and LTS $\mathcal{S}$ over $\mathcal{A}$, the LTS $\langle \mathcal{S} | \eta | \delta \rangle$ preserves $I^{-1}(P)$ with respect to $\mathcal{S}$.*

Theorems 3–4 do not show that our definition of PEP soundness is ‘minimal.’ Under Theorem 3, there could in theory exist *weaker* notions of PEP soundness that would still ensure that any three-way composition is sound with respect to the PDP’s property and PEP’s trace mediator. Similarly, there could exist weaker notions of PEP transparency that still guarantee preservation. The following theorems show such weaker definitions cannot exist:

Theorem 5. *Assume that $|\mathcal{E}| \geq 2$. An PEP η over $(\mathcal{A}, \mathcal{E})$ is sound with respect to I iff for any property P , for any edit automaton δ over $\mathcal{E}$ sound with respect to P , and for any LTS $\mathcal{S}$ over $\mathcal{A}$, the LTS $\langle \mathcal{S} | \eta | \delta \rangle$ complies with $I^{-1}(P)$.*

Theorem 6. *Assume that $|\mathcal{A}| \geq 2$. A PEP η over $(\mathcal{A}, \mathcal{E})$ is transparent iff for any property P , for any edit automaton δ transparent with respect to P , and for any LTS $\mathcal{S}$, the LTS $\langle \mathcal{S} | \eta | \delta \rangle$ preserves $I^{-1}(P)$ with respect to $\mathcal{S}$.*

4 Enforcement of Downward-Closed Properties

In the previous section, we have described how PEPs that are sound with respect to a trace mediator I can be used to ensure system compliance with $I^{-1}(P)$ when composing the system with a sound PDP for P . In many practical instances, however, PEPs which are not generally sound can still ensure compliance with more restricted classes of properties. For example, for the property P_{Consent} above, a PEP that suppresses *too many* Use events can still guarantee compliance. In this case, what allows compliant behavior despite an unsound PEP is that the property P_{Consent} is *downward-closed* with respect to the relation $\preceq_{\text{Use-}}$ such that $\sigma \preceq_{\text{Use-}} \sigma'$ iff σ is obtained from σ' by removing Use events.

An ‘overapproximating’ PEP such as the above has at least two potential advantages: because it produces some events less often, it may induce less runtime overhead; furthermore, its implementation might be easier to audit for soundness, since some constraints (e.g., the fact that Use events should be suppressed only on the PDP’s request) may be lifted. In this section, we study this overapproximation by considering the class of all properties $P \subseteq \mathcal{E}^*$ that are downward-closed with respect to some relation fixed $\preceq$ on $\mathcal{E}^*$:

Definition 14 (Downward Closure). Let $\preceq$ be a relation on $\mathcal{E}^*$. A property $P \subseteq \mathcal{E}^*$ is downward-closed under $\preceq$ iff $\forall \sigma, \sigma'. \sigma' \in P \wedge \sigma \preceq \sigma' \implies \sigma \in P$.

The following theorem follows straightforwardly from this definition:

Theorem 7. Let I and I' be two trace mediators such that $\forall \rho. I'(\rho) \preceq I(\rho)$. If a system S complies with $I(\rho)$, it complies with $I'(\rho)$. If S preserves $I'(\rho)$ with respect to some S_* , it preserves $I(\rho)$ with respect to S_* .

In the above example, one consequence of this theorem is that we have a PEP η that is sound with respect to some I' such that I' produces *fewer* **Consent** events than I , then any $\langle S|\eta|\delta_{\text{Consent}} \rangle$ satisfies $I^{-1}(P_{\text{Consent}})$. This is because the property P_{Consent} is downward-closed with respect to the relation $\preceq_{\text{Consent}+}$ such that $\sigma \preceq_{\text{Consent}+} \sigma'$ iff σ is obtained from σ' by adding **Consent** events.

Next, we consider the following notion of $\preceq$ -soundness for PEPs:

Definition 15. A PEP η is $\preceq$ -sound with respect to I iff whenever $s_\eta^0 \xrightarrow{\sigma \triangleright \rho}$ $s'_\eta \xrightarrow{a' \triangleright e'}_{\eta, i} s''_\eta \xrightarrow{e \triangleright \alpha}_{\eta, e} s'''_\eta$, then $\circ \sigma \cdot e$ is a prefix of some σ' such that $\sigma' \preceq I(\circ \rho \cdot \circ \alpha)$.

The next theorem formalizes the overapproximation discussed above in the case of **Use**: if a PEP η generates sequences of actions that are smaller (under $\preceq$) than the sequences of events edited by the PDP, then compliance is guaranteed.

Theorem 8. Let η be a PEP over $(\mathcal{A}, \mathcal{E})$ and δ an edit automaton over $\mathcal{E}$ sound with respect to P . If that η is $\preceq$ -sound with respect to I and P is downward-closed under $\preceq$, then $\langle S|\eta|\delta \rangle$ complies with $I^{-1}(P)$.

5 Practical Enforcement with Batches and Capabilities

The general model that we have described in Sections 3–4 must be instantiated for specific classes of events, PDPs, and downward-closed relations to be used with real-world tools. In this section, we specialize our theory to enforcement mechanisms that work with sets of simultaneous events (‘batches’, sometimes also called *timepoints* [6] in the literature) and distinguish between events that the PDP can cause or suppress [14]. This will allow us to implement our framework with the ENFGUARD enforcer [17] in Section 6. We first define batch edit automata, batch PEPs, and their properties, and introduce a relation $\preceq_{\text{mono}}$ whose downward closure contains properties that cannot be violated by generating ‘more’ or ‘fewer’ of certain events (Section 5.1). Then, we give a concrete batch PEP algorithm and prove sufficient ($\preceq$ -)soundness conditions (Section 5.2).

5.1 Batch Edit Automata and Batch PEPs

In this section, we consider the case where each $a \in \mathcal{A}$ and $e \in \mathcal{E}$ is a finite set, i.e., $\mathcal{A} \triangleq \mathcal{P}(A)_f$ and $\mathcal{E} \triangleq \mathcal{P}(E)_f$, where $\mathcal{P}(\cdot)_f$ denotes the set of all finite subsets. Accordingly, unlike in the previous sections, we now refer to elements of

$\mathcal{A}$ and $\mathcal{E}$ as *action batches* and *event batches*, and reserve the words *actions* and *events* for elements of A and E , respectively. As in ENFGUARD [17], we assume that the edit automaton serving as a PDP cannot insert or delete event batches, but only edit the *content* of the sets, and that which events can be caused or suppressed is known in advance: there is a set $C \subseteq E$ of *causable* events and a set $S \subseteq E$ of *suppressable* events. We call such automata *batch automata*; the pair (C, S) is called the *capabilities* of the automaton.

We first specialize our definitions of edit automata and PEPs:

Definition 16. A batch edit automaton (BEA) with capabilities (C, S) is an EA with transitions $s_\delta \xrightarrow{\ell \triangleright \langle \ell' \rangle} s'_\delta$ or $s_\delta \xrightarrow{\ell \triangleright \tau} s'_\delta$ only, where $\ell' - \ell \subseteq C$ and $\ell - \ell' \subseteq S$.

Definition 17. A batch PEP (BPEP) with capabilities (C, S) is an input-enabled PEP with transitions $s_\eta \xrightarrow{a \triangleright \langle e \rangle}_{\eta, i} s'_\eta$ and $s'_\eta \xrightarrow{e \triangleright \langle a \rangle}_{\eta, i} s''_\eta$ only where, whenever $s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s_\eta \xrightarrow{a' \triangleright \langle e' \rangle}_{\eta, i} s'_\eta \xrightarrow{e \triangleright \langle a \rangle}_{\eta, e} s''_\eta$, there exists $e'' \in \mathcal{E}$ and s'''_η such that $s_\eta \xrightarrow{a \triangleright \langle e'' \rangle} s'''_\eta$, $e'' - e' \subseteq C$, and $e' - e'' \subseteq S$.

Next, we define classes of properties that are particularly common in practice, namely those in which the addition or subtraction of a given event e to any batch of a compliant trace can never yield a non-compliant trace. If a property P can never be violated by just adding e events to an already compliant trace, we call it *e-monotonic*; when it can never be violated by just removing e events from such a trace, we call it *e-antimonotonic*.

Definition 18. A property $P \subseteq \mathcal{E}^*$ is *monotonic with respect to $e \in E$* iff for all $\sigma, \sigma' \in \mathcal{E}^*$ such that for all $i \in \mathbb{N}$, $\sigma'_i \in \{\sigma_i, \sigma_i \cup \{e\}\}$, $\sigma \in P \implies \sigma' \in P$.

Similarly, a property $P \subseteq \mathcal{E}^*$ is *antimonotonic with respect to $e \in E$* iff for all $\sigma, \sigma' \in \mathcal{E}^*$ such that for all $i \in \mathbb{N}$, $\sigma'_i \in \{\sigma_i, \sigma_i - \{e\}\}$, $\sigma \in P \implies \sigma' \in P$.

For the rest of this section, we fix two sets $M^+, M^- \subseteq E$ of events and consider the following relation $\preceq_{\text{mono}}$:

$$\sigma \preceq_{\text{mono}} \sigma' \iff |\sigma| = |\sigma'| \wedge \forall i \in \{1, \dots, |\sigma|\}. \sigma'_i - \sigma_i \subseteq M^+ \wedge \sigma_i - \sigma'_i \subseteq M^-.$$

The following proposition directly follows from Definition 18:

Theorem 9. A property P is downward-closed with respect to $\preceq_{\text{mono}}$ iff it is monotonic in every $e \in M^+$ and antimonotonic in every $e \in M^-$.

5.2 A Batch PEP Algorithm

In this subsection, we present a concrete batch PEP algorithm and give sufficient conditions for it to be (i) sound and transparent or at least (ii) $\preceq_{\text{mono}}$ -sound.

The algorithm. Our algorithm is defined via two functions $i_\eta : \mathcal{A}^* \times \mathcal{A} \rightarrow \mathcal{E}$ and $e_\eta : \mathcal{A}^* \times \mathcal{A} \times \mathcal{E} \rightarrow \mathcal{A}$ such that $s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s_\eta \xrightarrow{a' \triangleright \langle e' \rangle}_{\eta, i} s'_\eta$ iff $e' = i_\eta(\circ \rho, a')$ and

$s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s_\eta \xrightarrow{a' \triangleright \langle e' \rangle}_{\eta, i} s'_\eta \xrightarrow{e \triangleright \langle a \rangle}_{\eta, e} s''_\eta$ iff $a = \mathbf{e}_\eta(\circ \rho, a', e)$. We call the resulting PEP $\mathcal{B}(\mathbf{i}_\eta, \mathbf{e}_\eta)$. The function $\mathbf{i}_\eta$, called the *instrumentation mapping*, can be freely chosen. It maps an action to a sequence of events which is also allowed to depend on the history of past actions. The function $\mathbf{e}_\eta$, called the *enforcement mapping*, is defined in terms of the instrumentation mapping as well as three *handlers*: the causation handler $\mathbf{h}_C$, the preservation handler $\mathbf{h}_K$ ('K' stands for 'keep'), and the suppression handler $\mathbf{h}_S$. Each handler is the partial map describing how the enforcer should handle the PDP's edits: when the original set of actions is a' and the PDP causes some set of events e , the handler generates the set of actions $\mathbf{h}_C(e, a')$; when the PDP does not modify some events e , the handler generates the actions $\mathbf{h}_K(e, a')$; and when it suppresses some events e , the handler generates the actions $\mathbf{h}_S(e, a')$. Since the first argument of each handler is a set of events, there may be several ways to use the same handler to cause, preserve, or suppress the same events. For instance, to cause $\{e_1, e_2, e_3\}$, one can either call $\mathbf{h}_C(\{e_1\}, a')$, $\mathbf{h}_C(\{e_2\}, a')$, and $\mathbf{h}_C(\{e_3\}, a')$ consecutively and collect the generated actions, call $\mathbf{h}_C(\{e_1, e_2\}, a')$ and $\mathbf{h}_C(\{e_2, e_3\}, a')$, call only $\mathbf{h}_C(\{e_1, e_2, e_3\})$, etc. For the enforcer to be non-blocking, the domain of $\mathbf{h}_C$ (resp., $\mathbf{h}_K$, $\mathbf{h}_S$) must generate C (resp., E , S):

Definition 19. A set X is a generator of another set Y iff for any $y \in X$, there exists a finite subset $X' \subseteq X$ such that $y = \bigcup_{x \in X'} x$.

In the case where several decompositions of the set of events into a union of elements of the domain exists, we can use any function $\text{choose}_{Y;X} : Y \rightarrow \mathcal{P}(X)_f$ such that $\bigcup \text{choose}_{X;Y}(y) = y$ to select a valid decomposition. Given such a choice function, the handlers provide a convenient way for developers to define (sound) PEP reactions to arbitrary PDP edits, without having to explicitly define the PEP's reaction to every combination of edits.

Algorithm 1 gives the pseudocode of $\mathbf{e}_\eta$. First, the enforcer computes the events e' corresponding to the original actions and compares them to the edited events e returned by the PDP. If the PDP did not modify the events, the enforcer returns the original actions a' . Otherwise, it decomposes the sets of caused actions $e' - e$, preserved actions $e \cap e'$, and suppressed actions $e - e'$ over $\text{dom}(\mathbf{h}_C)$, $\text{dom}(\mathbf{h}_K)$, and $\text{dom}(\mathbf{h}_S)$ respectively, and calls the corresponding handlers, returning the union of all generated actions.

Correctness. The following theorem provides a set of conditions that together guarantee that $\mathcal{B}(\mathbf{i}_\eta, \mathbf{e}_\eta)$ is a sound and transparent PEP with capabilities (C, S) . We first state the theorem, and then give an intuitive interpretation of each of the conditions in terms of instrumentation auditing:

Theorem 10. Let $\mathbf{i}_\eta : \mathcal{A}^* \times \mathcal{A} \rightarrow \mathcal{E}$, $\mathbf{h}_C : \mathcal{P}(C)_f \rightarrow \mathcal{A} \rightarrow \mathcal{A}$, $\mathbf{h}_S : \mathcal{P}(S)_f \rightarrow \mathcal{A} \rightarrow \mathcal{A}$, $\mathbf{h}_K : \mathcal{P}(E)_f \rightarrow \mathcal{A} \rightarrow \mathcal{A}$ be given. Assume that (1) $\text{dom}(\mathbf{h}_C)$ is a generator of $\mathcal{P}(C)_f$ under $\cup$, (2) $\text{dom}(\mathbf{h}_K)$ is a generator of $\mathcal{P}(E)_f$ under $\cup$, (3) $\text{dom}(\mathbf{h}_S)$ is a generator of $\mathcal{P}(S)_f$ under $\cup$, (4) $\forall \rho, a, b. \mathbf{i}_\eta(\rho, a) \cup \mathbf{i}_\eta(\rho, b) = \mathbf{i}_\eta(\rho, a \cup b)$, where we set $\tau \cup s \triangleq s \cup \tau \triangleq s$, (5) $\forall \rho, c, a. \mathbf{i}_\eta(\rho, \mathbf{h}_C(c, a)) = c$, (6) $\forall \rho, k, a. \mathbf{i}_\eta(\rho, \mathbf{h}_K(k, a)) = k$, (7) $\forall \rho, s, a. \mathbf{i}_\eta(\rho, \mathbf{h}_S(s, a)) = \emptyset$, and (8) $\forall \rho, a. \mathbf{i}_\eta(\rho, a) = (I(\rho \cdot \langle a \rangle))_{|\rho|+1}$. Then

Algorithm 1 Batch PEP Algorithm

$$\begin{aligned}
 \mathbf{e}_\eta(\rho, a', e) = & \text{let } e' = \mathbf{i}_\eta(\rho, a') \text{ in} \\
 & (\text{if } e' = e \text{ then } a \\
 & \text{else let } c_1, \dots, c_{n_C} = \text{choose}_{\text{dom}(\mathbf{h}_C); C}(e - e'); \\
 & \quad k_1, \dots, k_{n_K} = \text{choose}_{\text{dom}(\mathbf{h}_K); K}(e \cap e'); \\
 & \quad s_1, \dots, s_{n_S} = \text{choose}_{\text{dom}(\mathbf{h}_S); S}(e' - e) \\
 & \text{in } \bigcup_{i=1}^{n_C} \mathbf{h}_C(c_i, a') \cup \bigcup_{i=1}^{n_K} \mathbf{h}_K(k_i, a') \cup \bigcup_{i=1}^{n_S} \mathbf{h}_S(s_i, a'))
 \end{aligned}$$

$\mathcal{B}(\mathbf{i}_\eta, \mathbf{e}_\eta)$ is a batch PEP with capabilities (C, S) that is sound and transparent with respect to I .

This theorem provides a checklist (1)–(8) that can be used to audit the correctness of the system’s instrumentation when our PEP algorithm is used together with a sound PDP. The checks to be performed are as follows:

- (1–3) Is the causation handler (resp. preservation, suppression handler) defined for all causable events (resp. for all events, for all suppressable events)?
- (4) Does the instrumentation mapping map single actions to events independently, i.e., does a set of n actions map to the same events as the union of the events that each of the n actions maps to?
- (5–6) When the causation (resp. preservation) handler receives events, does it generate actions that map exactly (through $\mathbf{i}_\eta$) to the events it received?
- (7) Does the suppression handler only generate actions that map to $\emptyset$?
- (8) Does the instrumentation mapping $\mathbf{i}_\eta$ implement exactly I ?

For all events $e \in \mathcal{E}$ such that $e \in \mathbf{i}_\eta(\rho, \{x\}) \implies \mathbf{i}_\eta(\rho, \{x\}) = \{e\}$ (i.e., events that are never generated together with another event), note that the part of the preservation handler related to e can be straightforwardly defined as $a' \mapsto \mathbf{h}_K(\{e\}, a') = \{x \in a' \mid e \in \mathbf{i}_\eta(\rho, \{x\})\}$, triggering all actions in a' that map to e .

In practice, auditing these requirements, especially (6) and (8), can be challenging, as it requires checking correct instrumentation and enforcement for *all possible events*: the instrumentation must *always* provide *exactly* the right events *whenever necessary*; the enforcer must *always* generate *exactly* the right events.

For $\preceq_{\text{mono}}$ soundness, we can significantly weaken conditions (5)–(8):

Theorem 11. *Let $\mathbf{i}_\eta$, $\mathbf{h}_C$, $\mathbf{h}_S$, and $\mathbf{h}_K$ be given. Assume that (1)–(4) and (7) are as in Theorem 10, (5) $\forall c, a. \mathbf{i}_\eta(\rho, \mathbf{h}_C(c, a)) - c \subseteq C \cap M^+ \wedge c - \mathbf{i}_\eta(\rho, \mathbf{h}_C(c, a)) \subseteq M^-$, (6) $\forall \rho, k, a. \mathbf{i}_\eta(\rho, \mathbf{h}_K(k, a)) - k \subseteq C \cap M^+ \wedge k - \mathbf{i}_\eta(\rho, \mathbf{h}_K(k, a)) \subseteq S \cap M^-$, (8) $\forall \rho, a. \mathbf{i}_\eta(\rho, a) \preceq (I(\rho \cdot \langle a \rangle))_{|\rho|+1}$. Then $\mathcal{B}(\mathbf{i}_\eta, \mathbf{e}_\eta)$ is $\preceq_{\text{mono}}$ -sound with respect to I .*

The questions to be answered for (5), (6), and (8) are now the following:

- (5–6) a. When a handler receives an event that is *not both* antimonotonic and suppressable, does it always generate an action that maps to this event?
- b. When a handler generates an action that maps to an event that is *not both* monotonic and causable, did it always receive this event originally?

Actions	Event	Sets
Any reading or writing of some of user u 's data for purpose p	$\text{Use}(u, p)$	S, M^-
Any user input from u giving consent for purpose p	$\text{Consent}(u, p)$	M^+
Any user input from u revoking consent for purpose p	$\text{Revoke}(u, p)$	M^-
Any user input from u containing a deletion request	$\text{Request}(u)$	M^-
Any call to a function that delete all of user u 's data	$\text{Delete}(u)$	C, M^+

Table 1: Specification of the trace mediator I in Minitwitter

- (8) a. Does the implementation mapping i_η ensure that every event that is *not* monotonic is always logged when an action mapping to it occurs?
 b. Does the implementation mapping i_η ensure that every event that is *not* antimonotonic is only ever logged when an action mapping to it occurs?

6 Case study

INSTRLIB. We have implemented the batch PEP algorithm in Section 5 in a Python library called INSTRLIB. The library consists of roughly 1,500 lines of Python code with two instrumentation layers: a low-level layer, which allows for instrumenting arbitrary Python function calls and attributes, and a high-level layer providing off-the-shelf enforcement hooks for Django web applications. The latter uses a fixed set of actions (**read**, **write**, **input**, **output**, **execute**) to capture database reads and writes, user inputs and outputs, and calls to class members. The library allows developers to specify the PEP logic as in Algorithm 1 by providing an instrumentation mapping and handlers. Its architecture is shown in Appendix B. It has bindings to the state-of-the-art ENFGUARD tool [17].

Minitwitter. We now showcase the usage of INSTRLIB by instrumenting a micro-blogging app for compliance with two privacy requirements. The target app has 435 lines of Python code and allows users to view their and other users' timeline, follow other users, and post short messages. Additionally, the app shows one of two advertisement messages on the user's timeline depending on the content of their posts. It also displays a privacy banner that prompts the user to accept or reject the use of their data for marketing purposes. We are interested in enforcing the two following requirements: (1) whenever data is used for marketing purposes, the user has given (and not revoked) consent; (2) the user can request the suppression of all of their data within one minute. Here, the events of interest are $\text{Use}(u, p)$, denoting "user u 's data is used for purpose p ;" $\text{Consent}(u, p)$ (resp. $\text{Revoke}(u, p)$), denoting "user u gives (resp. revokes) consent to use their data for purpose p ;" $\text{Request}(u)$, denoting "user u requests deletion of their data;" and $\text{Delete}(u)$, denoting "user u 's data is deleted." The PDP is assumed to be able to cause **Delete** and suppress **Use**. The trace mediator I , which is part of the software specification, is shown in Table 1. As in most practical instances, this trace mediator is informal.

View	Baseline ($\propto n$)					Instrumented ($\propto n$)				
	10^2	10^3	10^4	10^5	10^6	10^2	10^3	10^4	10^5	10^6
timeline	54	54	58	58	66	56 +2	58 +4	69 +11	70 +12	75 +9
post	64	62	63	62	61	67 +4	68 +6	66 +3	67 +5	67 +6
consent	47	46	47	47	47	53 +6	54 +8	55 +8	54 +7	54 +7
request	—	—	—	—	—	58	59	59	58	72

Table 2: Runtime latency (ms) over 20 repetitions

To instrument Minitwitter, developers proceed in three steps. *First*, they provide the event signature and the policy of interest in ENFGUARD’s [17] specification logic: Metric First-Order Temporal Logic (MFOTL) [9]. Here, the policy is

$$\begin{aligned} & \Box(\forall u. (\text{Use}(u, \text{"marketing"}) \rightarrow (\neg \text{Revoke}(u, \text{"marketing"}) \text{ S } \text{Consent}(u, \text{"marketing"}))) \\ & \wedge (\text{Request}(u) \rightarrow \Diamond_{[0,60]} \text{Delete}(u))). \end{aligned}$$

Second, the developers use the built-in Django bindings for INSTRLIB to associate actions to database reads and writes (actions **read** and **write**), user inputs to views (action **input**), and function calls (action **execute**). Functions with a specific purpose of processing can be marked with that purpose (here, **"marketing"**). At runtime, based on the current call stack, INSTRLIB injects the current purposes of processing into **read** and **write** actions. This step requires about 40 lines of code in the Python files describing the app’s models and URLs.

Third, the developers describe the instrumentation mapping and handlers. This requires about 50 lines of code in a single Python file. The instrumentation mapping maps database **read** or **writes** to **Use** events, consent banner clicks (captured by specific **input** actions) to either **Consent** or **Revoke**, clicks on a special ‘Delete My Data’ button (captured by other **input** actions) to **Request**, and executions of a special function **delete_data(u)** that erases all of a user’s data (captured by an **execute** action) to **Delete**. Two handlers are implemented: a suppression handler for **Use** that returns **None** instead of the actual content of object fields and prevents their overwriting; and a causation handler for **Delete** that calls **delete_data**. By default, INSTRLIB provides simple preservation handlers as described in Section 5.2, which are sufficient when i_η maps each action to at most one event. All enforcement-related code is showed in Appendix A.

In Table 2, we report the latency of four of Minitwitter’s views with and without instrumentation with INSTRLIB: viewing the **timeline**, **posting** a message, giving **consent**, and **requesting** deletion of one’s data, for different values of the number n of posts in the database. The runtime overhead is <15 ms per request.

Auditing Minitwitter’s implementation. In the policy above, **Consent** and **Delete** are monotonic, whereas **Request**, **Revoke**, and **Use** are antimonotonic. We can now go through the checklist provided by Theorem 11. For (1–3), we check the existence of a causation handler for **Delete**, a suppression handler for **Use**, and preservation handlers for all events. Regarding preservation handlers, we note that, as described above, INSTRLIB’s default implementation provides simple preservation handler that are sufficient with our choice of i_η —this also allows

us to check (6). Condition (4) is implemented by INSTRLIB by design. For the Delete causation handler, we answer (5a) positively by checking that the handler does call the `delete_data` function, whose behavior matches the informal description in Table 1. Condition (5b) is vacuous since Delete is monotonic and causable. Condition (7) is trivially fulfilled since our suppression handlers return `None`. For (8a), we must check that Request, Use, and Revoke events are always emitted when actions that map to them according to Table 1 occur. Inspecting the interface of the application, we control that the buttons for revocation of consent and deletion requests map to the views whose `inputs` we have instrumented. Similarly, we check that all fields that contain personal data emit `read` and `write` events and that all functions performing marketing are marked as such. Finally, for (8b), we must check that Consent and Delete are only logged when the corresponding system actions as described in Table 1 happen. To this end, we control that Consent is only generated by the `input` corresponding to clicking ‘yes’ in the banner and, similarly, that the Delete event is only generated by the `execute` action of function `delete_data`.

7 Conclusions and Future Work

To the best of our knowledge, we have provided the first formal account of instrumentation in runtime enforcement. Our framework provides necessary and sufficient conditions for the correctness of the composition of a system, PDP, and PEP and is implemented in the INSTRLIB library.

Future work includes extending our auditing methodology to validate the implementation of runtime enforcement mechanisms in large applications with complex specifications and further optimizing INSTRLIB.

References

1. Luca Aceto, Ian Cassar, Adrian Francalanza, and Anna Ingólfssdóttir. On Runtime Enforcement via Suppressions. In Sven Schewe and Lijun Zhang, editors, *29th International Conference on Concurrency Theory (CONCUR 2018)*, Leibniz International Proceedings in Informatics (LIPIcs), pages 34:1–34:17, Dagstuhl, Germany. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
2. Luca Aceto, Ian Cassar, Adrian Francalanza, and Anna Ingólfssdóttir. Bidirectional runtime enforcement of first-order branching-time properties. *Logical Methods in Computer Science*, 19, 2023.
3. Rajeev Alur and David L Dill. A theory of timed automata. *Theoretical computer science*, 126(2):183–235, 1994.
4. Guangdong Bai, Liang Gu, Tao Feng, Yao Guo, and Xiangqun Chen. Context-aware usage control for android. In *Security and Privacy in Communication Networks: 6th International ICST Conference, SecureComm 2010, Singapore, September 7-9, 2010. Proceedings 6*, pages 326–343. Springer, 2010.
5. David Basin, Germano Caronni, Sarah Ereth, Matúš Harvan, Felix Klaedtke, and Heiko Mantel. Scalable offline monitoring of temporal specifications. *Formal Methods in System Design*, 49:75–108, 2016.

6. David Basin, Felix Klaedtke, Samuel Müller, and Birgit Pfitzmann. Runtime monitoring of metric first-order temporal properties. In *IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science (2008)*, pages 49–60. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2008.
7. Ian Cassar, Adrian Francalanza, Luca Aceto, and Anna Ingólfssdóttir. Developing theoretical foundations for runtime enforcement. *CoRR*, abs/1804.08917, 2018.
8. Hadil Charafeddine, Khalil El-Harake, Yliès Falcone, and Mohamad Jaber. Runtime enforcement for component-based systems. In *Proceedings of the 30th annual ACM symposium on applied computing*, pages 1789–1796, 2015.
9. Jan Chomicki. Efficient checking of temporal integrity constraints using bounded history encoding. *ACM Transactions on Database Systems (TODS)*, 20(2):149–186, 1995.
10. Úlfar Erlingsson. *The inlined reference monitor approach to security policy enforcement*. PhD thesis, 2004.
11. Úlfar Erlingsson and Fred B Schneider. Sasi enforcement of security policies: A retrospective. In *Proceedings of the 1999 workshop on New security paradigms*, pages 87–95, 1999.
12. Yliès Falcone and Mohamad Jaber. Fully automated runtime enforcement of component-based systems with formal and sound recovery. *International Journal on Software Tools for Technology Transfer*, 19(3):341–365, 2017.
13. Simon Godik, eds. Tim Moses, and OASIS XACML Technical Committee. extensible access control markup language (xacml) version 1.0. OASIS Standard oasis-xacml-1.0, OASIS, February 2003. OASIS Standard, 18 February 2003.
14. François Hublet, David Basin, and Srđan Krstić. Real-time policy enforcement with metric first-order temporal logic. In *European Symposium on Research in Computer Security*, pages 211–232. Springer, 2022.
15. François Hublet, David Basin, and Srđan Krstić. Enforcing the GDPR. In *European Symposium on Research in Computer Security*, pages 400–422. Springer, 2023.
16. François Hublet, Leonardo Lima, David Basin, Srđan Krstić, and Dmitriy Traytel. Proactive real-time first-order enforcement. In *International Conference on Computer Aided Verification*, pages 156–181. Springer, 2024.
17. François Hublet, Leonardo Lima, David Basin, Srđan Krstić, and Dmitriy Traytel. Scaling up proactive enforcement. In *International Conference on Computer Aided Verification*, 2025.
18. Jay Ligatti, Lujo Bauer, and David Walker. Edit automata: Enforcement mechanisms for run-time security policies. *International Journal of Information Security*, 4:2–16, 2005.
19. Jay Ligatti and Srikar Reddy. A theory of runtime enforcement, with results. In *Computer Security–ESORICS 2010: 15th European Symposium on Research in Computer Security, Athens, Greece, September 20–22, 2010. Proceedings 15*, pages 87–100. Springer, 2010.
20. Nancy A. Lynch and Mark R. Tuttle. *An introduction to input/output automata*. Laboratory for Computer Science, Massachusetts Institute of Technology, 1988.
21. Srinivas Pinisetty, Ylies Falcone, Thierry Jéron, Hervé Marchand, Antoine Rollet, and Omer Landry Nguena Timo. Runtime enforcement of timed properties. In *International Conference on Runtime Verification*, pages 229–244. Springer, 2012.
22. Siegfried Rasthofer, Steven Arzt, Enrico Lovat, and Eric Bodden. Droidforce: Enforcing complex, data-centric, system-wide policies in android. In *2014 Ninth International Conference on Availability, Reliability and Security*, pages 40–49. IEEE, 2014.

23. Samira Sarraf. Optus breach occurred due to a coding error, alleges acma. *CSO Online*, 2022.
24. Fred Schneider. Enforceable security policies. *ACM Trans. Inf. Syst. Secur.*, 3(1):30–50, 2000.

A Enforcement code for Minitwitter

```

1  ...
2
3  from instrlib.django.orm import InstrumentORM
4
5  from .enforcer import logger
6
7  def info_user(user):
8      try:
9          return str(user)
10     except:
11         return ""
12
13  @InstrumentORM(logger,
14                  {"User.first_name", "User.last_name", "User.is_staff", "User.
15                   email", "User.save", "User.last_login"},
16                  info = info_user, events = {'read', 'write', 'execute'})
17  class User(AbstractUser):
18
19      def delete_data(self):
20          self.twit.all().delete()
21          self.follower.all().delete()
22          self.following.all().delete()
23
24  def info_follow(follow):
25      try:
26          return str(object.__getattribute__(follow, 'follower'))
27      except:
28          return ""
29
30  @InstrumentORM(logger,
31                  {"Follow.follower", "Follow.following", "Follow.date", "Follow
32                   "},
33                  info = info_follow, events = {'read', 'write'})
34  class Follow(CommonInfo):
35      ...
36
37  def info_twit(twit):
38      try:
39          return str(object.__getattribute__(twit, 'author'))
40      except:
41          return ""
42
43  @InstrumentORM(logger,
44                  {"Twit.content", "Twit.posted_on", "Twit.updated_on", "Twit.
45                   author", "Twit.save"},
46                  info = info_twit, events = {'read', 'write'})
47  class Twit(CommonInfo):
48      ...

```

models.py

```

1  ...
2
3  from instrlib.django.url import InstrumentURL
4  from twitt.enforcer import logger

```

```

5
6 urlpatterns = [
7     path('admin/', admin.site.urls),
8     path('', include('twitt.urls')),
9 ]
10
11 to_instrument = {
12     "twitt.views.SetCookieConsentView",
13     "twitt.views.DeleteAllView",
14 }
15
16 urlpatterns = InstrumentURL(logger, to_instrument, events = {'input'})(
    urlpatterns)

```

urls.py

```

1 class HomeView(LoginRequiredMixin, TemplateView):
2     template_name = 'index.html'
3
4     @with_purpose('marketing')
5     def generate_advertisement(self):
6         ...

```

views.py

```

1 from typing import Any
2
3 from instrlib.event import Event
4 from instrlib.pdp import EnfGuard
5 from instrlib.logger import Logger
6 from instrlib.pep import InstrumentationMapping, PEP
7 from instrlib.schema import Schema
8
9 from Twitter.settings import INSTRLIB_EXE, INSTRLIB_FORMULA, INSTRLIB_LOG,
    INSTRLIB_SIG
10
11 # Handlers
12
13 def none_handler(event_name, event_args, response, *args, **kwargs):
14     return None
15
16 def delete_handler(events):
17     from twitt.models import User
18
19     for event in events:
20
21         user_name = event['args'][0]
22         try:
23             user = User.objects.get(username=user_name)
24         except User.DoesNotExist:
25             return
26
27         user.delete_data()
28
29 # Schema
30
31 schema = Schema()
32 schema.add('Use', [str, str])
33 schema.add('Consent', [str, str])
34 schema.add('Request', [str])
35 schema.add('Delete', [str])
36
37 # PDP
38
39 pdp = EnfGuard(INSTRLIB_EXE, INSTRLIB_SIG, INSTRLIB_FORMULA, log_file =
    INSTRLIB_LOG)

```

```

40
41 # PEP
42
43 suppression_handlers : dict[str | tuple[str, ...], Any] = {
44     ('Use') : none_handler,
45 }
46 causation_handlers : dict[str | tuple[str, ...], Any] = {
47     ('Delete') : delete_handler,
48 }
49
50 def read_mapping(action):
51     return Event('Use', action.args[4], action.args[5])
52
53 def write_mapping(action):
54     return Event('Use', action.args[5], action.args[6])
55
56 def input_mapping(action):
57     if action.args[:3] == ('SetCookieConsentView', 'accept', 'true'):
58         return Event('Consent', action.args[3], 'marketing')
59     elif action.args[:3] == ('DeleteAllView', 'delete', 'true'):
60         return Event('Request', action.args[3])
61     else:
62         return None
63
64 def execute_mapping(action):
65     if action.args[:2] == ('User', 'delete_data'):
66         return Event('Delete', action.args[3])
67     else:
68         return None
69
70 instrumentation_mapping = InstrumentationMapping({
71     'read': read_mapping,
72     'write': write_mapping,
73     'input': input_mapping,
74     'execute': execute_mapping,
75 })
76
77 pep = PEP(
78     suppression_handlers = suppression_handlers,
79     causation_handlers = causation_handlers,
80     instrumentation_mapping = instrumentation_mapping
81 )
82
83 # Logger
84
85 logger = Logger(pep, schema, pdp)

```

enforcer.py

B Architecture of Instrlib

The architecture of INSTRLIB, shown in Figure 4, involves six different kinds of parallel threads: worker threads; writer and reader threads; proactive worker threads; the PDP; and a timer thread. These threads interact according to two main flows. The first, *reactive* [16] flow (in green in Figure 4) starts with an instrumented function being run by a worker serving a user request; the events to be logged are computed and placed into a queue (1) that is read by the writer thread (2). The writer thread passes the events to the PDP (3) and places the original worker request in another queue (4). The reader thread reads the outputs of the PDP (5) and realigns them with the content of the second queue




Fig. 4: Architecture of INSTRLIB

(6) before passing it back to the worker (7). Additionally, the reader thread can start *proactive* workers (8) to perform causation. The second, *proactive* [16] flow (in red) is initiated by a clock tick. It proceeds as before (2–6) but can only perform causation (8). This second flow is used when, e.g., an action has to be performed before a deadline irrespective of the presence of user inputs.

A Proofs for Section 3 (The Extended Enforcement Model)

Let $\sqsubseteq$ denote prefix inclusion on traces, i.e., $\sigma \sqsubseteq \sigma' \triangleq \exists \sigma''. \sigma' = \sigma \cdot \sigma''$.

Theorem 3. *An edit automaton δ over $\mathcal{E}$ is sound with respect to P iff for any PEP η over $(\mathcal{A}, \mathcal{E})$ that is sound with respect to I and LTS $\mathcal{S}$ over $\mathcal{A}$, the LTS $\langle \mathcal{S}|\eta|\delta \rangle$ complies with $I^{-1}(P) \triangleq \{\rho \in \text{dom}(I) \mid I(\rho) \in P\}$.*

Proof. $\implies$ Assume that δ is sound with respect to P , η is sound, and let ρ be a trace of $\langle \mathcal{S}|\eta|\delta \rangle$, i.e. $(s_{\mathcal{S}}^0, s_{\eta}^0, s_{\delta}^0, \emptyset) \xrightarrow{\rho} (s_{\mathcal{S}}^n, s_{\eta}^n, s_{\delta}^n, b)$. By Definition 11, there exist σ and σ' such that $s_{\delta}^0 \xrightarrow{\sigma' \triangleright \sigma} s_{\delta}^n$ and $s_{\eta}^0 \xrightarrow{\sigma \triangleright \rho} s_{\eta}^n$, where $\sigma = \circ \sigma$ and $\rho = \circ \rho$. Expanding the definition of δ 's soundness, we get $\text{pre}(\sigma) \subseteq P$. By Definition 15, since η is sound, we get that, for any $i \leq |\rho|$, $I(\circ \rho_{..i}) \sqsubseteq \circ \sigma_{..i}$. In particular, $I(\rho) \in \text{pre}(\sigma) \subseteq P$. We conclude that $\rho \in I^{-1}(P)$.

$\Leftarrow$ Assume that δ is not sound with respect to P , i.e., there exists an execution $s_{\delta}^0 \xrightarrow{\sigma' \triangleright \sigma} s_{\delta}$ of δ such that $\text{pre}(\sigma) \not\subseteq P$. Let $i \leq |\sigma|$ such that $\sigma_{..i} \notin P$. Fix $\mathcal{E} = \mathcal{A}$. Consider the identity PEP $\text{id}_{\mathcal{A}}$ which (soundly) maps every action to itself with respect to $I = \text{id}$ and let $\mathcal{S} = \mathcal{U}_{\mathcal{A}}$ be the non-deterministic ‘universal’ system with state space $\{\perp\}$ that can take any transition in $\mathcal{A}$ at any time. Then, by Definition 11, there exists an execution $s_{\langle \mathcal{S}|\text{id}_{\mathcal{A}}|\delta \rangle}^0 \xrightarrow{\sigma_{..i}} s$ (note that the three-way composition performs at most one non- τ action per transition). Since $\sigma_{..i} \notin P$, then $\langle \mathcal{S}|\text{id}_{\mathcal{A}}|\delta \rangle$ does not comply with $P = I^{-1}(P)$.

Theorem 4. *An edit automaton δ over $\mathcal{E}$ is transparent with respect to P iff for any transparent PEP η over $(\mathcal{A}, \mathcal{E})$ and LTS $\mathcal{S}$ over $\mathcal{A}$, the LTS $\langle \mathcal{S}|\eta|\delta \rangle$ preserves $I^{-1}(P)$ with respect to $\mathcal{S}$.*

Proof. $\implies$ Assume that δ is transparent with respect to P , η is transparent, and let $\rho = (a_i)_{i \in \mathbb{N}}$ be a trace of $\mathcal{S}$ (i.e., $s^0 \xrightarrow{\rho}$) such that $\rho \in I^{-1}(P)$. Choose $\sigma = I(\rho) \in P$. By our choice of ρ , we have $\sigma \in P$. Expanding the definition of δ 's transparency, we obtain an execution $s_{\delta}^0 \xrightarrow{\sigma \triangleright \sigma}$ where $\sigma = \langle \langle \sigma_i \rangle \rangle_{i \in \mathbb{N}}$. Using Definition 13, we can finally construct an execution $s_{\langle \mathcal{S}|\eta|\delta \rangle}^0 \xrightarrow{\rho}$ where $\rho = \langle \langle \rho_i \rangle \rangle_{i \in \mathbb{N}}$.

$\Leftarrow$ Assume that δ is not transparent with respect to P , i.e., there exists an execution $s_{\delta}^0 \xrightarrow{\sigma' \triangleright \sigma}$ of δ such that $\sigma' \in P$ but $\sigma \notin P$. Consider the identity PEP $\text{id}_{\mathcal{A}}$, which is transparent, and let $I = \text{id}$ and $\mathcal{S} = \mathcal{U}_{\mathcal{A}}$. Then $s^0 \xrightarrow{\sigma'}$ and $s_{\langle \mathcal{S}|\text{id}_{\mathcal{A}}|\delta \rangle}^0 \xrightarrow{\sigma}$, but $\sigma \notin P$, hence $\langle \mathcal{S}|\text{id}_{\mathcal{A}}|\delta \rangle$ does not preserve $I^{-1}(P)$ with respect to $\mathcal{S}$.

Theorem 5. *Assume that $|\mathcal{E}| \geq 2$. An PEP η over $(\mathcal{A}, \mathcal{E})$ is sound with respect to I iff for any property P , for any edit automaton δ over $\mathcal{E}$ sound with respect to P , and for any LTS $\mathcal{S}$ over $\mathcal{A}$, the LTS $\langle \mathcal{S}|\eta|\delta \rangle$ complies with $I^{-1}(P)$.*

Proof. $\implies$ Identical to $(\implies)$ in the proof of Theorem 3.

$\Leftarrow$ Assume that η is not sound. Let (σ, ρ) of minimal length such that $s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s'_\eta, s'_\eta \xrightarrow{a' \triangleright e'} s''_\eta, s''_\eta \xrightarrow{e \triangleright \alpha} s'''_\eta, \circ \rho \cdot \circ \alpha \in \text{dom}(I)$, and $\sigma^\dagger \triangleq I(\circ \rho \cdot \circ \alpha)$ is not a prefix of $\circ \sigma \cdot e$.

Let $P \triangleq \{\sigma' \in \mathcal{E}^* \mid \exists \sigma'' \in \mathcal{E}^*. \sigma' = \circ \sigma \cdot e \cdot \sigma'' \vee \circ \sigma \cdot e = \sigma' \cdot \sigma''\}$ and $\mathcal{S} = \mathcal{U}_A$. Clearly, the property P is not empty. Consider an edit automaton δ that enforces P by first editing the sequence of events to be exactly $\circ \sigma$; then, when passed e' , rewriting it to e ; and later not modifying the trace anymore. By Definition 11, there exists an execution $s_{\langle \mathcal{S} | \eta | \delta \rangle}^0 \xrightarrow{\circ \rho} (\emptyset, s'_\delta, s'_\eta, \varepsilon)$ for some s'_δ . There are two cases:

- If $\circ \sigma \cdot e$ is not a (proper) prefix of $\sigma^\dagger$, consider an execution of $\langle \mathcal{S} | \eta | \delta \rangle$ that performs one step from $(\emptyset, s'_\delta, s'_\eta, \varepsilon)$, generating $\sigma' = \circ \rho \cdot \circ \alpha$. Then $\sigma' \notin P$.
- If $\circ \sigma \cdot e$ is a proper prefix of $\sigma^\dagger$, let $f_1 \in \mathcal{E}$ and $f_2 \in \mathcal{E}^*$ such that $\sigma^\dagger = \circ \sigma \cdot e \cdot \langle f_1 \rangle \cdot f_2$. Further, let $f_3 \in \mathcal{E} - \{f_1\}$. Modify P to $P' \triangleq P - \{\sigma' \in \mathbb{T}_\mathcal{E} \mid \exists \sigma''. \sigma' = \sigma \cdot f \cdot \sigma''\}$. Consider an edit automaton δ' that soundly enforces P' by first editing the sequence of events to be exactly $\circ \sigma$; then, when passed e' , rewriting it to e ; and later preventing reaching any suffix of $\circ \sigma \cdot f$ by editing f to $e \cdot \langle f_3 \rangle \cdot f_2$ if necessary. Again, consider an execution of $\langle \mathcal{S} | \eta | \delta \rangle$ that performs one step from $(\emptyset, s'_\delta, s'_\eta, \varepsilon)$, generating $\sigma' = \circ \rho \cdot \circ \alpha$. Then $\sigma' \notin P'$.

In both cases, we exhibited a system $\mathcal{S}$ and sound edit automaton δ for some property P such that $\langle \mathcal{S} | \eta | \delta \rangle$ does not comply with $I^{-1}(P)$.

Theorem 6. Assume that $|\mathcal{A}| \geq 2$. A PEP η over $(\mathcal{A}, \mathcal{E})$ is transparent iff for any property P , for any edit automaton δ transparent with respect to P , and for any LTS $\mathcal{S}$, the LTS $\langle \mathcal{S} | \eta | \delta \rangle$ preserves $I^{-1}(P)$ with respect to $\mathcal{S}$.

Proof. $\Rightarrow$ Identical to $(\Rightarrow)$ in the proof of Theorem 4.

$\Leftarrow$ If the PEP η is not non-blocking, then any execution when η blocks on a execution of $\mathcal{U}_A$ under the trivial policy $P = \mathcal{E}^*$ shows that the enforced system does not preserve $I^{-1}(P)$.

Now, assume that there exist $s_\eta^0 \xrightarrow{\sigma \triangleright \rho} s'_\eta, s'_\eta \xrightarrow{a' \triangleright e'} s''_\eta$, and $s''_\eta \xrightarrow{e \triangleright \alpha} s'''_\eta$ such that $\alpha \neq \langle a' \rangle$. Let $\mathcal{S}$ be a system with state $\mathbb{S} \subseteq \mathcal{A}^*$ that can only produce a single trace $\rho \cdot \langle a' \rangle \cdot \rho'$ for some $\rho' \in \mathcal{A}^\omega$, keeping the trace prefix it has seen so far in its state. Let δ be a trivial edit automaton for property $P = \mathcal{E}^\omega$. Consider an execution of $\langle \mathcal{S} | \eta | \delta \rangle$ that continues from $(\rho \cdot \alpha, \emptyset, s'''_\eta, \varepsilon)$. As in the proof of Theorem 5, there are three cases:

- If none of α and $\langle a' \rangle$ is a prefix of the other, then the system is blocked after s'''_η , being forced by the PEP to diverge from $\rho \cdot \langle a' \rangle \cdot \rho'$.
- If α is a proper prefix of $\langle a' \rangle$, i.e., $\alpha = \varepsilon$, then $s'''_\eta = s'_\eta$ by Definition 9 and the system loops, failing to generate any infinite trace.
- If $\langle a' \rangle$ is a proper prefix of α , i.e., $\alpha = \langle a' \rangle \cdot \mathbf{b}$ for some $\mathbf{b} \in \mathcal{A}^* - \{\varepsilon\}$, then observe that since $|\mathcal{A}| \geq 2$, in the above, we can pick $\rho' \notin \langle a' \rangle \cdot \mathbf{b} \cdot \mathcal{A}^*$. Then no execution of $\langle \mathcal{S} | \eta | \delta \rangle$ that continues from $(\rho \cdot \alpha, \emptyset, s'''_\eta, \varepsilon)$ can preserve the original trace of the system.

B Proofs for Section 4 (Enforcement of Downward-Closed Properties)

Theorem 8. *Let η be a PEP over $(\mathcal{A}, \mathcal{E})$ and δ an edit automaton over $\mathcal{E}$ sound with respect to P . If that η is $\preceq$ -sound with respect to I and P is downward-closed under $\preceq$, then $\langle \mathcal{S}|\eta|\delta \rangle$ complies with $I^{-1}(P)$.*

Proof. Analogously to Theorem 3 ($\implies$), we prove that $\langle \mathcal{S}|\eta|\delta \rangle$ complies with $I^{-1}(P)$ using our downward-closure assumption. Let ρ be a trace of $\langle \mathcal{S}|\eta|\delta \rangle$. Since $\langle \mathcal{S}|\eta|\delta \rangle$ complies with $I^{-1}(P)$, then $\rho \in I^{-1}(P)$, i.e., $\sigma \triangleq I(\rho) \in P$. By our first assumption, $I^*(\rho) \preceq I(\rho)$. By our second assumption applied to σ and $\sigma' \triangleq I(\rho)$, we get $I^*(\rho) \in P$, i.e., $\rho \in I^{*-1}(P)$. Hence, $\langle \mathcal{S}|\eta|\delta \rangle$ complies with $I^{*-1}(P)$.

Theorem 12. *Let η be a PEP over $(\mathcal{A}, \mathcal{E})$ and δ an edit automaton over $\mathcal{E}$. Assume that $\langle \mathcal{S}|\eta|\delta \rangle$ preserves $I^{-1}(P)$ with respect to $\mathcal{S}$. If $\forall \rho. I(\rho) \preceq I^*(\rho)$ and P is downward-closed under $\preceq$, then $\langle \mathcal{S}|\eta|\delta \rangle$ preserves $I^{*-1}(P)$ with respect to $\mathcal{S}$.*

Proof. Let $I^{*-1}(P)$ be a trace of $\langle \mathcal{S}|\eta|\delta \rangle$. By our first assumption, $I(\rho) \preceq I^*(\rho)$. By our second assumption, $I(\rho) \in P$, i.e., $\rho \in I^{-1}(P)$. Since $\langle \mathcal{S}|\eta|\delta \rangle$ preserves $I^{-1}(P)$ with respect to $\mathcal{S}$, then ρ is the trace of an execution of $\mathcal{S}$. Hence, $\langle \mathcal{S}|\eta|\delta \rangle$ preserves $I^{*-1}(P)$ with respect to $\mathcal{S}$.

C Proofs for Section 5 (Practical Enforcement with Batches and Capabilities)

Theorem 10. *Let $i_\eta : \mathcal{A}^* \times \mathcal{A} \rightarrow \mathcal{E}$, $h_C : \mathcal{P}(C)_f \rightarrow \mathcal{A} \rightarrow \mathcal{A}$, $h_S : \mathcal{P}(S)_f \rightarrow \mathcal{A} \rightarrow \mathcal{A}$ be given. Assume that (1) $\text{dom}(h_C)$ is a generator of $\mathcal{P}(C)_f$ under $\cup$, (2) $\text{dom}(h_K)$ is a generator of $\mathcal{P}(E)_f$ under $\cup$, (3) $\text{dom}(h_S)$ is a generator of $\mathcal{P}(S)_f$ under $\cup$, (4) $\forall \rho, a, b. i_\eta(\rho, a) \cup i_\eta(\rho, b) = i_\eta(\rho, a \cup b)$, where we set $\tau \cup s \triangleq s \cup \tau \triangleq s$, (5) $\forall \rho, c, a. i_\eta(\rho, h_C(c, a)) = c$, (6) $\forall \rho, k, a. i_\eta(\rho, h_K(k, a)) = k$, (7) $\forall \rho, s, a. i_\eta(\rho, h_S(s, a)) = \emptyset$, and (8) $\forall \rho, a. i_\eta(\rho, a) = (I(\rho \cdot \langle a \rangle))_{|\rho|+1}$. Then $\mathcal{B}(i_\eta, e_\eta)$ is a batch PEP with capabilities (C, S) that is sound and transparent with respect to I .*

Proof. Since (i) $\text{dom}(h_C)$ is a generator of $\mathcal{P}(E)_f$ under $\cup$ and (ii) $\text{dom}(h_S)$ is a generator of $\mathcal{P}(S)_f$, then the function e_η is total. Transparency is straightforward. For soundness, we compute

$$\begin{aligned}
& i_\eta(a_C \cup a_K \cup a_S) \\
&= i_\eta \left(\bigcup_{i=1}^{n_C} h_C(c_i, a') \cup \bigcup_{i=1}^{n_K} h_K(k_i, a') \cup \bigcup_{i=1}^{n_S} h_S(s_i, a') \right) \\
&= \left(\bigcup_{i=1}^{n_C} i_\eta(h_C(c_i, a')) \right) \cup \left(\bigcup_{i=1}^{n_K} i_\eta(h_K(k_i, a')) \right) \cup \left(\bigcup_{i=1}^{n_S} i_\eta(h_S(s_i, a')) \right) \quad \text{by (iii)} \\
&= \left(\bigcup_{i=1}^{n_C} c_i \right) \cup \left(\bigcup_{i=1}^{n_K} k_i \right) \cup \left(\bigcup_{i=1}^{n_S} \emptyset \right) \quad \text{by (iv)-(v)}
\end{aligned}$$

$$= e - e' \cup e \cap e' = e.$$

by def. $\mathbf{e}_\eta$

Theorem 11. *Let $\mathbf{i}_\eta$, $\mathbf{h}_C$, $\mathbf{h}_S$, and $\mathbf{h}_K$ be given. Assume that (1)–(4) and (7) are as in Theorem 10, (5) $\forall c, a. \mathbf{i}_\eta(\rho, \mathbf{h}_C(c, a)) - c \subseteq C \cap M^+ \wedge c - \mathbf{i}_\eta(\rho, \mathbf{h}_C(c, a)) \subseteq M^-$, (6) $\forall \rho, k, a. \mathbf{i}_\eta(\rho, \mathbf{h}_K(k, a)) - k \subseteq C \cap M^+ \wedge k - \mathbf{i}_\eta(\rho, \mathbf{h}_K(k, a)) \subseteq S \cap M^-$, (8) $\forall \rho, a. \mathbf{i}_\eta(\rho, a) \preceq (I(\rho \cdot \langle a \rangle))_{|\rho|+1}$. Then $\mathcal{B}(\mathbf{i}_\eta, \mathbf{e}_\eta)$ is $\preceq_{\text{mono}}$ -sound with respect to I .*

Proof. By Theorems 8 and 10.